

### 1. What are the advantages of a CNN over a fully connected DNN for image classification?

Answer:---->Convolutional Neural Networks (CNNs) have several advantages over fully connected Deep Neural Networks (DNNs) for image classification. Some of the main advantages include:

**Spatial Awareness:** CNNs have a built-in notion of spatial awareness. They take into account the spatial relationships between the pixels of an image, while fully connected DNNs treat an image as a 1D array of pixels.

**Parameter Efficiency:** CNNs use a lot fewer parameters than fully connected DNNs for image classification, which results in faster training times and a reduction in overfitting.

**Translation Invariance:** CNNs have the ability to detect features at multiple positions in an image, which makes them translation-invariant.

**Local Connectivity:** In a CNN, each neuron receives input only from a small region of the previous layer, which reduces the number of parameters and the computation required.

**Sharing Weights:** In a CNN, the same filters are applied to multiple regions of an image, which reduces the number of parameters and helps prevent overfitting.

These advantages make CNNs particularly well-suited for image classification tasks, where the spatial relationships between pixels are important.

### 2. Consider a CNN composed of three convolutional layers, each with $3 \times 3$ kernels, a stride of 2, and "same" padding. The lowest layer outputs 100 feature maps, the middle one outputs 200, and the top one outputs 400. The input images are RGB images of $200 \times 300$ pixels.

What is the total number of parameters in the CNN? If we are using 32-bit floats, at least how much RAM will this network require when making a prediction for a single instance? What about when training on a mini-batch of 50 images?

Answer:---->The number of parameters in a convolutional layer depends on the number of filters (also called feature maps), the size of the filters, and the number of input channels. If we have  $3 \times 3$  filters, a stride of 2, and "same" padding, then the number of parameters in each layer can be calculated as follows:

Number of parameters in the first layer =  $(3 * 3 * 3 + 1) * 100 = 2800 + 100 = 2900$  parameters  
 Number of parameters in the second layer =  $(3 * 3 * 100 + 1) * 200 = 18000 + 200 = 18200$  parameters  
 Number of parameters in the third layer =  $(3 * 3 * 200 + 1) * 400 = 72000 + 400 = 72400$  parameters  
 So the total number of parameters in the CNN is  $2900 + 18200 + 72400 = 82500$  parameters.

When making a prediction for a single instance, the RAM requirement will depend on the size of the activations and gradients, which are computed during forward and backward propagation, respectively. Assuming that each activation and gradient value is a 32-bit float, the total size of activations and gradients in the CNN can be estimated as follows:

The lowest layer outputs a feature map of size  $(200 / 2) \times (300 / 2) \times 100 = 100 \times 150 \times 100$   
 The middle layer outputs a feature map of size  $(100 / 2) \times (150 / 2) \times 200 = 50 \times 75 \times 200$   
 The top layer outputs a feature map of size  $(50 / 2) \times (75 / 2) \times 400 = 25 \times 38 \times 400$   
 So the total size of activations and gradients in the CNN is  $100 \times 150 \times 100 + 50 \times 75 \times 200 + 25 \times 38 \times 400 = 9000000 + 5625000 + 3750000 = 18375000$  float values.

The total size in RAM will be  $18375000 * 4 \text{ bytes/float} = 735 \text{ MB}$ .

When training on a mini-batch of 50 images, the RAM requirement will be multiplied by the batch size. The total size in RAM will be  $735 \text{ MB} * 50 = 36.75 \text{ GB}$

### 3. If your GPU runs out of memory while training a CNN, what are five things you could try to solve the problem?

Answer:---->1.Reduce batch size: Reducing the batch size can significantly reduce the GPU memory requirement and help alleviate the memory constraints.

2.Use mixed precision: Using mixed precision training, where the model parameters are stored in lower-precision data types, such as float16, can significantly reduce memory usage.

3.Decrease the complexity of the model: Decreasing the number of layers or filters in the model can help reduce memory consumption.

4.Gradient checkpointing: This technique enables you to train very deep models by only keeping a subset of activations in memory, and recomputing the rest when needed.

5. Transfer learning: Consider using pre-trained models, instead of training the entire model from scratch. This approach is significantly less memory-intensive and can help reduce the amount of memory required to store the model parameters.

#### 4. Why would you want to add a max pooling layer rather than a convolutional layer with the same stride?

Answer:--->Max pooling is a down-sampling operation that reduces the spatial dimensions of the input tensor while retaining the important features. By doing so, max pooling reduces the number of parameters in the model, which helps to prevent overfitting and reduces the computation required during training. Max pooling also has the benefit of being more robust to small translations and distortions in the input.

A convolutional layer with the same stride as max pooling would have the same effect on the spatial dimensions of the input, but would still require more parameters to be learned by the model. Additionally, a convolutional layer with a stride of 2 is essentially learning a down-sampling operation, which is less interpretable and harder to optimize than max pooling.

In general, max pooling is used in conjunction with convolutional layers as it helps to reduce the spatial dimensions and retain the important features of the input, while also simplifying the model and reducing the computation required during training.

#### 5. When would you want to add a local response normalization layer?

Answer:----->Local response normalization, also known as LRN, is a normalization technique used in Convolutional Neural Networks (CNNs). The main goal of LRN is to address the internal covariate shift problem, which is the phenomenon where the distribution of activations changes during the training process. This can cause optimization difficulties and slow convergence.

LRN performs normalization across the feature maps (depth-wise) within a mini-batch of instances, normalizing each activation value with respect to the mean and variance of its surrounding activations.

LRN is usually added to a CNN architecture when training on large image datasets, as it has been shown to improve the performance in some cases. However, LRN is not used very frequently in modern deep learning, as other normalization techniques, such as batch normalization, have proven to be more effective and efficient in most cases.

#### 6. Can you name the main innovations in AlexNet, compared to LeNet-5? What about the main innovations in GoogLeNet, ResNet, SENet, and Xception?

Answer:-----> Yes, here's a brief overview of the main innovations in each of the mentioned deep convolutional neural networks:

##### 1. AlexNet:

AlexNet introduced the concept of using deep convolutional neural networks for large-scale image classification, and was the winner of the 2012 ImageNet competition. It was one of the first deep convolutional neural networks to use rectified linear unit (ReLU) activation instead of the traditional sigmoid activation. AlexNet also introduced the idea of dropout, a regularization technique to reduce overfitting.

##### 2. GoogLeNet:

GoogLeNet (also known as InceptionNet) introduced the concept of "Inception modules" - a combination of multiple convolutional and pooling layers in a single module. This allowed the network to be more compact and computationally efficient while still preserving its ability to learn hierarchical representations of the input data. GoogLeNet was also the first deep convolutional neural network to use average pooling instead of fully connected layers for prediction.

##### 3. ResNet:

ResNet introduced the concept of "residual connections" - allowing the network to learn residual functions with reference to the input, instead of learning the complete mapping from input to output. This greatly improved the training of deep neural networks, allowing for much deeper architectures without the risk of vanishing gradients or poor performance.

##### 4. SENet:

SENet (Squeeze-and-Excitation networks) introduced the concept of using attention mechanisms in deep convolutional neural networks. The attention mechanism allows the network to focus on different parts of the input data depending on its importance, rather than applying the same processing to all parts. This leads to improved performance and a more interpretable model.

##### 5. Xception:

Xception introduced the concept of using depthwise separable convolutions - a computationally efficient form of convolution where the spatial convolution and channel-wise convolution are performed separately. This greatly reduced the number of parameters in the network, allowing for

### 7. What is a fully convolutional network? How can you convert a dense layer into a convolutional layer?

Answer:---->A fully convolutional network (FCN) is a type of convolutional neural network (CNN) where the fully connected dense layers in traditional CNNs are replaced by convolutional layers with large receptive fields. This allows the network to take inputs of variable size and produce outputs of the same size, which is useful for tasks like semantic segmentation.

To convert a dense layer into a convolutional layer, you can add a 1x1 convolutional layer. This effectively applies a dense connection to each of the spatial locations of the previous layer's output. This can be useful for incorporating learned global features into a high-resolution feature map.

### 8. What is the main technical difficulty of semantic segmentation?

Answer:----> The main technical difficulty of semantic segmentation is the challenge of accurately assigning a semantic label to each pixel in an image. This requires both the ability to precisely localize objects within the image and to accurately distinguish between object classes. The difficulty can be compounded by the presence of complex scenes with multiple objects, occlusions, and variations in scale, position, and orientation. To address these difficulties, it is often necessary to use complex deep neural network architectures, such as fully convolutional networks, and to make use of large annotated datasets for training.

### 9. Build your own CNN from scratch and try to achieve the highest possible accuracy on MNIST.

Answer:--->

### 10. Use transfer learning for large image classification, going through these steps:

- Create a training set containing at least 100 images per class. For example, you could classify your own pictures based on the location (beach, mountain, city, etc.), or alternatively you can use an existing dataset (e.g., from TensorFlow Datasets).
- Split it into a training set, a validation set, and a test set.
- Build the input pipeline, including the appropriate preprocessing operations, and optionally add data augmentation.
- Fine-tune a pretrained model on this dataset.

